

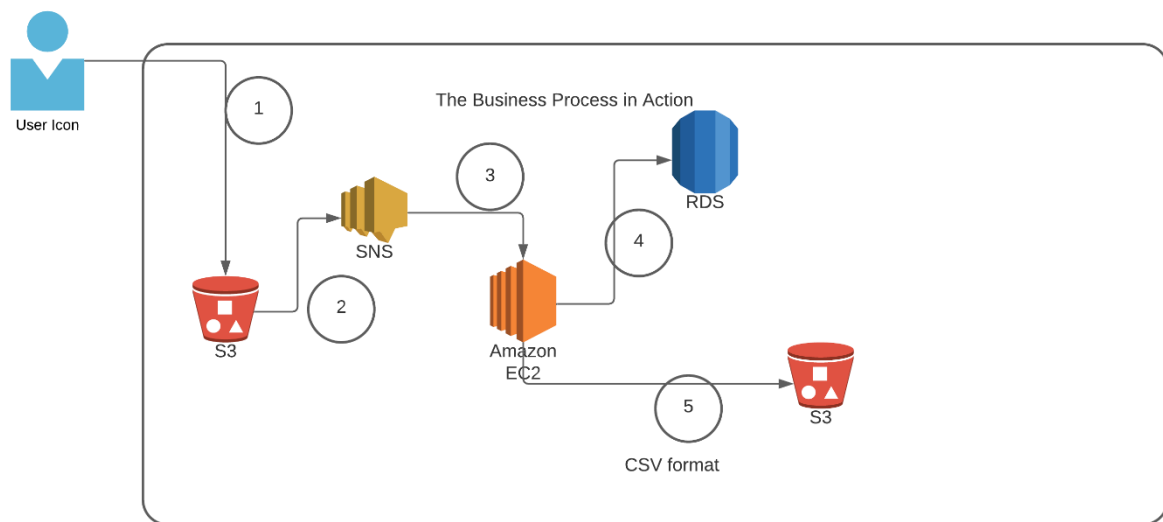
Declaration	
Questions in this exercise are intentionally complex and could be convoluted or confusing. This is by design and to simulate real life situations where customers seldom give crystal clear requirements and ask unambiguous questions.	

I have read the above statement and agree to these conditions	
I AGREE	Abhishek Pramod Agnihotri
	<Enter your name above this line to indicate that you are in agreement>

Instructions	
Every screenshot requested in this workbook is compulsory and carries 1 marks	
Your AWS account ID must be clearly visible in every screenshot using the AWS console; missing id or using someone else's id is not permitted. Such cases will be considered as plagiarism and severe penalty will be imposed.	
All screenshots must be in the order mentioned under "Expected Screenshots" for every step	
DO NOT WAIT UNTIL THE LAST MINUTE. The program office will not extend the project submission deadline under any circumstances.	
The file should be renamed in the format BATCH_FIRSTNAME_LASTNAME_PROJECT1. For example: PGPCCMAY18_VIJAY_DWIVEDI_PROJECT1.pdf	

Resource Clean Up	
Cloud is always pay per use model and all resources/services that we consume are chargeable. Cleaning up when you've completed your lab or project is always necessary. This is true whether you're doing a lab or implementing a project at your workplace.	
After completing the lab, make sure to delete each resource created in reverse chronological order. Each AWS Academy session lasts for 4 hours by default, although you can extend a session to run longer by pressing the start button to reset your session timer. At the end of each session, any resources you created in the account will be preserved. Some AWS resources, such as EC2 instances, may be automatically shut down, while other resources, such as RDS instances will be left running.	

Architecture diagram



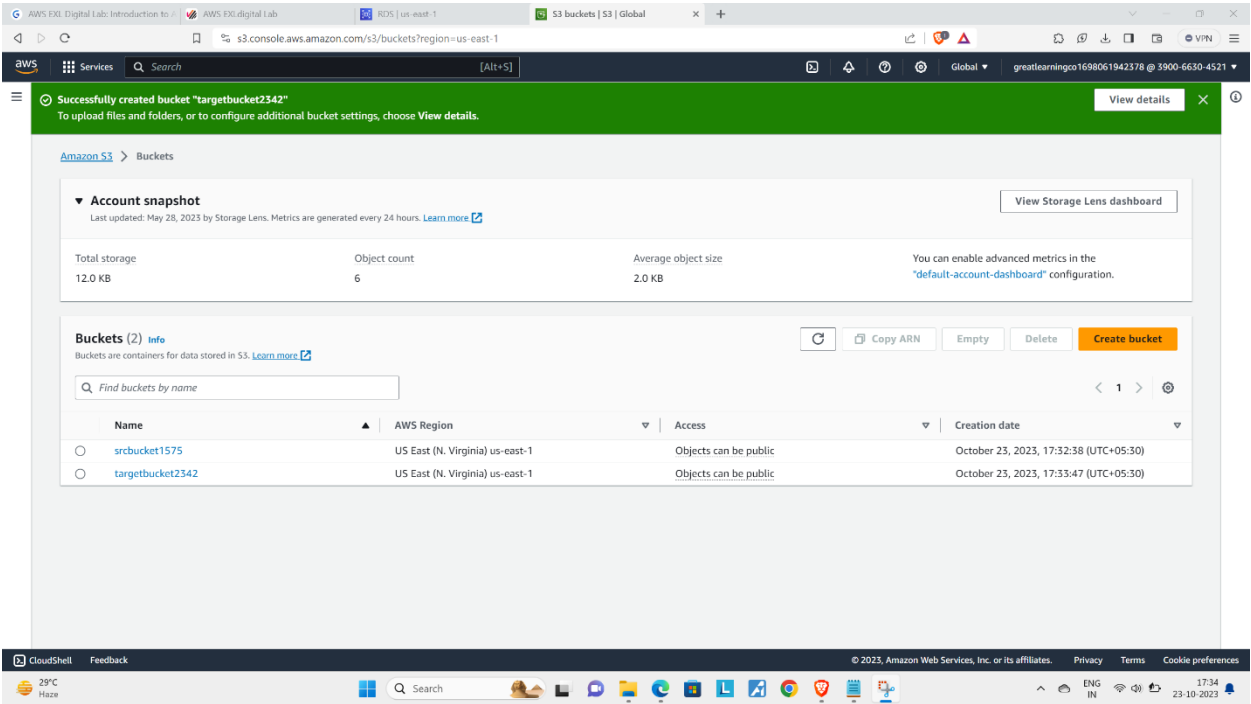
Architecture Implementation

1	The customer uploads the invoice data to S3 bucket in a text format as per their guidelines and policies. This bucket will have a policy to auto delete any content that is more than 1 day old (24 hours).
2	An event will trigger in the bucket that will place a message in SNS topic
3	A custom program running in EC2 will subscribe to the SNS topic and get the message placed by S3 event
4	The program will use S3 API to read from the bucket, parse the content of the file and create a CSV record and save the details in an RDS database
5	The program will use S3 API to write CSV record to destination S3 bucket as new S3 object.
Note	The custom program codebase and sample invoice have been shared along with this workbook on the LMS.

Step 1: SNS and S3 topic creation

Step number	a
Step name	Creation of Source and target buckets
Instructions	1) Navigate to S3 using the Services button at the top of the screen 2) Select "Create Bucket" 3) Enter a source bucket name and use the default options for the rest of the fields 4) Click on "Create Bucket" 5) Repeat the above steps to create a target bucket
Expected screenshots	1) Screen showing created S3 source and target buckets

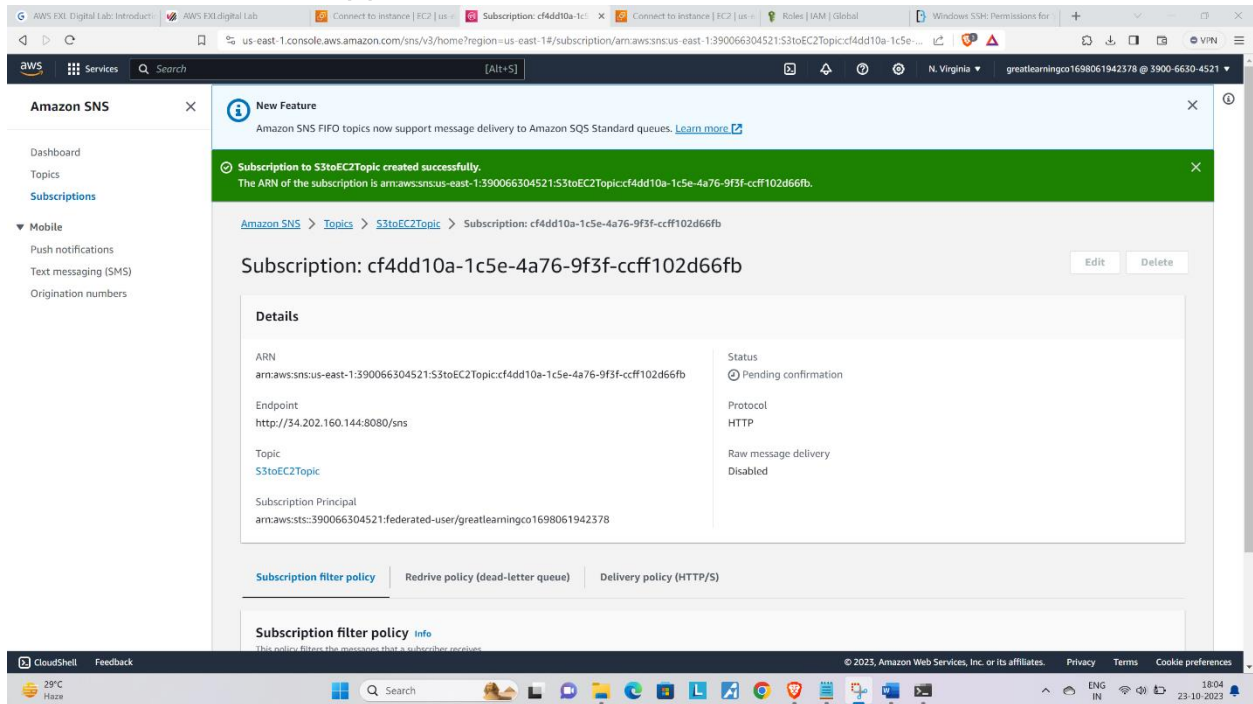
<Insert screenshot for a(1) here>



Step number	b
Step name	Creation of SNS subscription

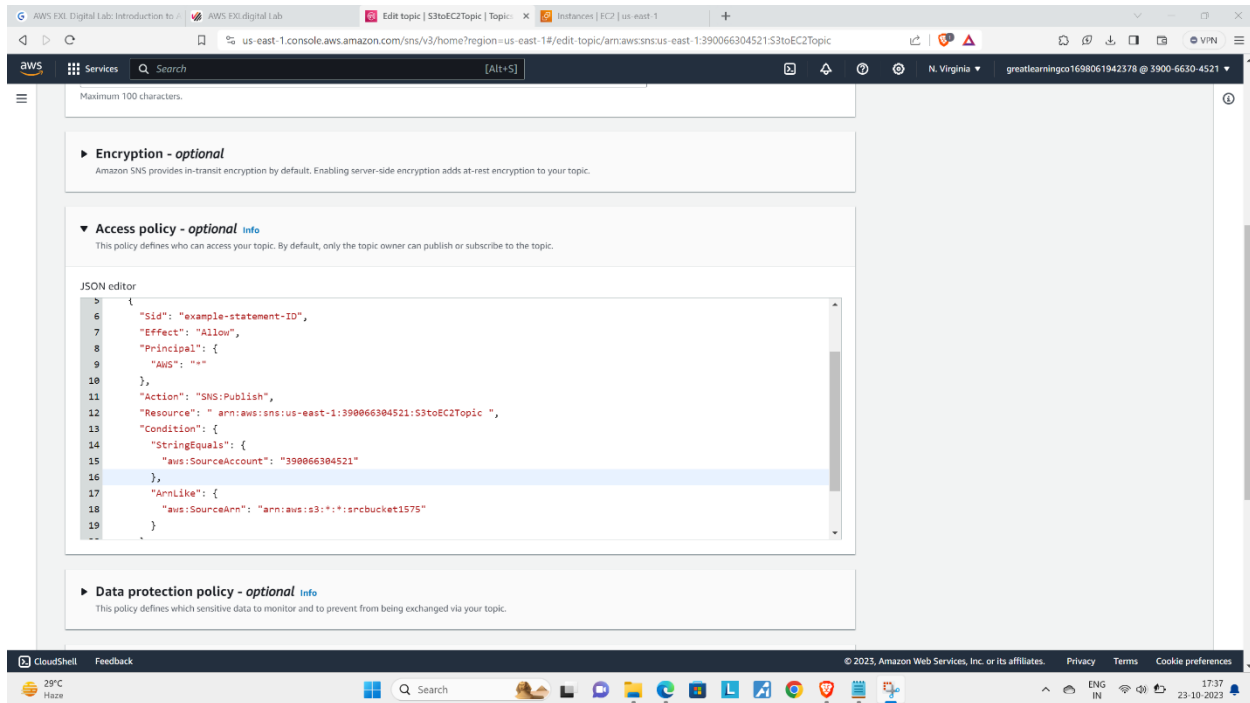
Instructions	<ol style="list-style-type: none"> 1) Navigate to SNS -> Topics 2) Click on "Create Topic" 3) Enter the following fields Name : S3toEC2Topic The other options can be ignored for now 4) Click on Create Topic
Expected screenshots	<ol style="list-style-type: none"> 1) Creation of SNS topic

<Insert screenshot for b(1) here>



Step number	c
Step name	Modification of SNS Access Policy
Instructions	<p>1) Navigate to SNS -> Topics and select the topic created in the previous step</p> <p>2) Note down the ARN shown in the topic details</p> <p>2) Click on Edit and select "Access Policy".</p> <p>3) Replace the text in the JSON editor with the following</p> <pre>{ "Version": "2012-10-17", "Id": "example-ID", "Statement": [{ "Sid": "example-statement-ID", "Effect": "Allow", "Principal": { "AWS": "*" }, "Action": ["SNS:Publish"], "Resource": "SNS-topic-ARN", "Condition": { "ArnLike": { "aws:SourceArn": "arn:aws:s3:*:*:bucket-name" }, "StringEquals": { "aws:SourceAccount": "bucket-owner-account-id" } } }] }</pre> <p>4) Replace the bold text with the SNS topic ARN, source bucket name and your AWS account ID respectively.</p> <p>5) Click on Save Changes</p>
Expected screenshots	1) JSON Editor screen

<Insert screenshot for c(1) here>



Step number	d
Step name	Configuring SNS notifications for S3
Instructions	<ol style="list-style-type: none"> 1) Navigate to S3 and select the source bucket created in Step 1 (a) 2) Select Properties and scroll down to Event Notifications and select it 3) Select "Create Event Notification" 4) Fillup the details as follows <ul style="list-style-type: none"> Name : S3PutEvent Select PUT from the list of radio buttons Destination : Select SNS Topic SNS : Select S3ToEC2Topic 5) Save Changes
Expected screenshots	1) Event Configuration Screen

<Insert screenshot for d(1) here>

The screenshot displays the AWS IAM console interface for the 'srcbucket1575' bucket. The 'Properties' tab is active, showing various configuration options. A prominent red error message states: 'You don't have permission to get AWS CloudTrail data events details. You or your AWS administrator must update your IAM permissions to allow cloudtrail:DescribeTrails. After you obtain the necessary permission, choose Refresh. Learn more about [identity and access management in Amazon S3](#).' Below this, the 'Event notifications' section lists one notification for 'S3PutEvent' with the destination 'S3toEC2Topic'. The 'Amazon EventBridge' section is currently set to 'Off', and the 'Transfer acceleration' section is also 'Off'. The top navigation bar shows the user is logged in as 'greatlearningco1698061942378'.

Step 2: Run the custom program in the EC2 instance

Step number	a
Step name	Creation of the EC2 instance and RDS instance
Instructions	<p>1) Navigate to EC2 -> Instances</p> <p>2) Create an EC2 instance with the following parameters</p> <p>AMI : Amazon Linux 2</p> <p>VPC : Default</p> <p>Security group : Ports 22 and 8080 should be opened</p> <p>3) Navigate to RDS</p> <p>4) Create an RDS instance with the following parameters:</p> <p>Engine type : MySql</p> <p>Template : Dev/Test</p> <p>Set the username and password as required</p> <p>DB Instance class : Burstable</p> <p>Instance type : t3.micro</p> <p>Storage type : General purpose SSD (gp2)</p> <p>Public Access : Yes</p> <p>VPC Security group : Create New ()</p> <p>Under Additional Configuration, add an initial database name. Take note of this name as it will be required later.</p> <p>Uncheck "Enable Enhanced Monitoring"</p> <p>Ensure that the security group created by the RDS deployment has port 3306 open for all incoming connections from all sources.</p>
Expected screenshots	<p>1) List of instances after creation of EC2 instance</p> <p>2) List of RDS instances</p>

<Insert screenshot for a(1) here>

<Insert screenshot for a(2) here>

us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#instances:

Instances (1/1) info

Find Instance by attribute or tag (case-sensitive)

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4
agninstance	i-01fe6ffd615e7a848	Running	t2.micro	2/2 checks passed	No alarms	us-east-1d	ec2-34-202-160-144.co...	34.202.160.1

Instance: i-01fe6ffd615e7a848 (agninstance)

Details Security Networking Storage Status checks Monitoring Tags

▼ Instance summary info

Instance ID	Public IPv4 address	Private IPv4 addresses
i-01fe6ffd615e7a848 (agninstance)	34.202.160.144 [open address]	172.31.25.184
IPv6 address	Instance state	Public IPv4 DNS
-	Running	ec2-34-202-160-144.compute-1.amazonaws.com [open address]
Hostname type	Private IP DNS name (IPv4 only)	Elastic IP addresses
IP name: ip-172-31-25-184.ec2.internal	ip-172-31-25-184.ec2.internal	
Answer private resource DNS name (IPv4)	Instance type	
	t3.micro	

Amazon RDS

Dashboard Databases Query Editor Performance insights Snapshots Exports in Amazon S3 Automated backups Reserved instances Proxies

Subnet groups Parameter groups Option groups Custom engine versions Zero-ETL integrations **New**

Events Event subscriptions

Recommendations 1 Certificate update 1

Databases (1)

Consider creating a Blue/Green Deployment to minimize downtime during upgrades. You may want to consider using Amazon RDS Blue/Green Deployments and minimize your downtime during upgrades. A Blue/Green Deployment provides a staging environment for changes to production databases. [RDS User Guide](#) [Aurora User Guide](#)

Group resources Modify Actions Restore from S3 Create database

Filter by databases

DB identifier	Status	Role	Engine	Region & AZ	Size	Actions	CPU	Current activity	Maintenance
database-1	Available	Instance	MySQL Community	us-east-1a	db.t3.micro	1 Action	3.47%	0 Connections	none

Step number b

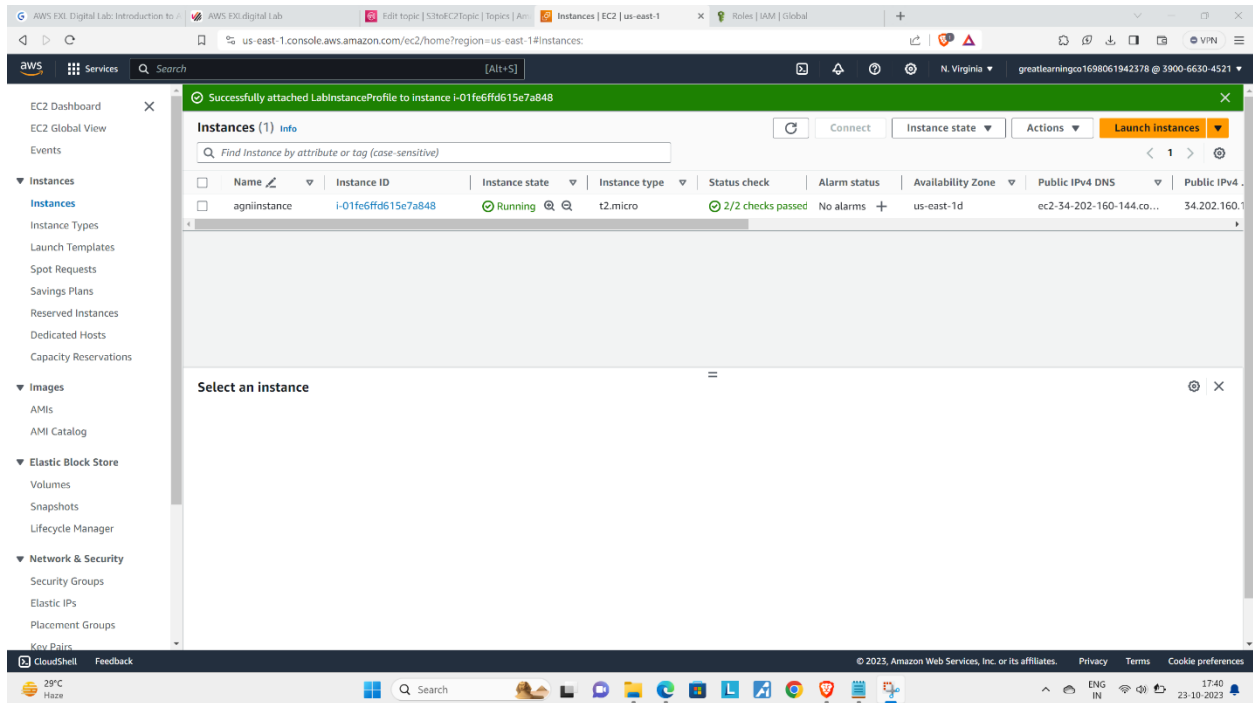
Step name Assignment of IAM role for EC2 instance

Instructions

- 1) Navigate back to EC2- > Instances
- 2) Select the EC2 instance created in the previous step and select Actions-> Security -> Modify IAM role
- 3) Select the role LabInstanceProfile from the dropdown and click on Save

Expected 1) Modify IAM role screen
screenshots

<Insert screenshot for b(1) here>



Step number	c
Step name	Configuration and Uploading of custom program
Instructions	<ol style="list-style-type: none"> 1) Download the file docproc-new.zip on your machine 2) Unzip the downloaded file 3) Enter the unzipped folder and open the file views.py in the API folder using a text editor 4) In line number 19-24, modify the target bucket name to the one created in Step 2 (a) and modify the hostname, username, password and database variables to the values set while creating the RDS database and save the file 5) Copy the folder docproc-new to the home folder of the EC2 instance created in Step 3(a) using scp. Use the command given below

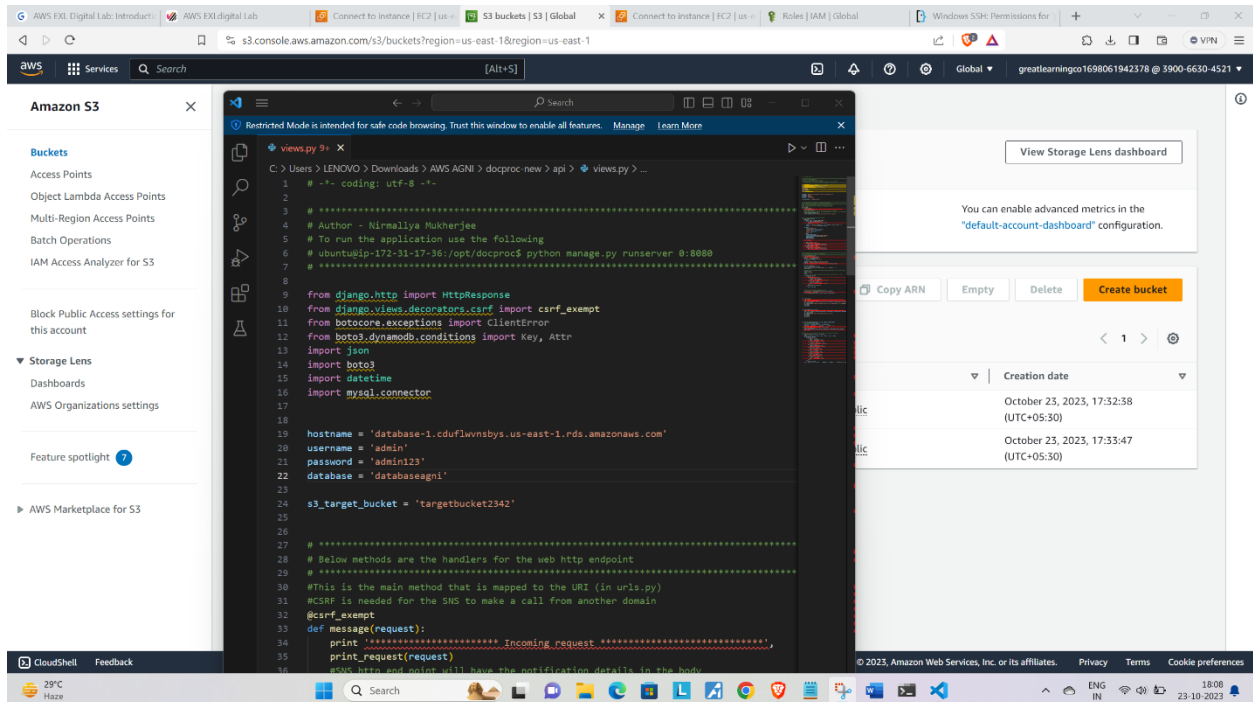
```
scp -i <pem> -r ./docproc-new ec2-  
user@<ip>:/home/ec2-user
```

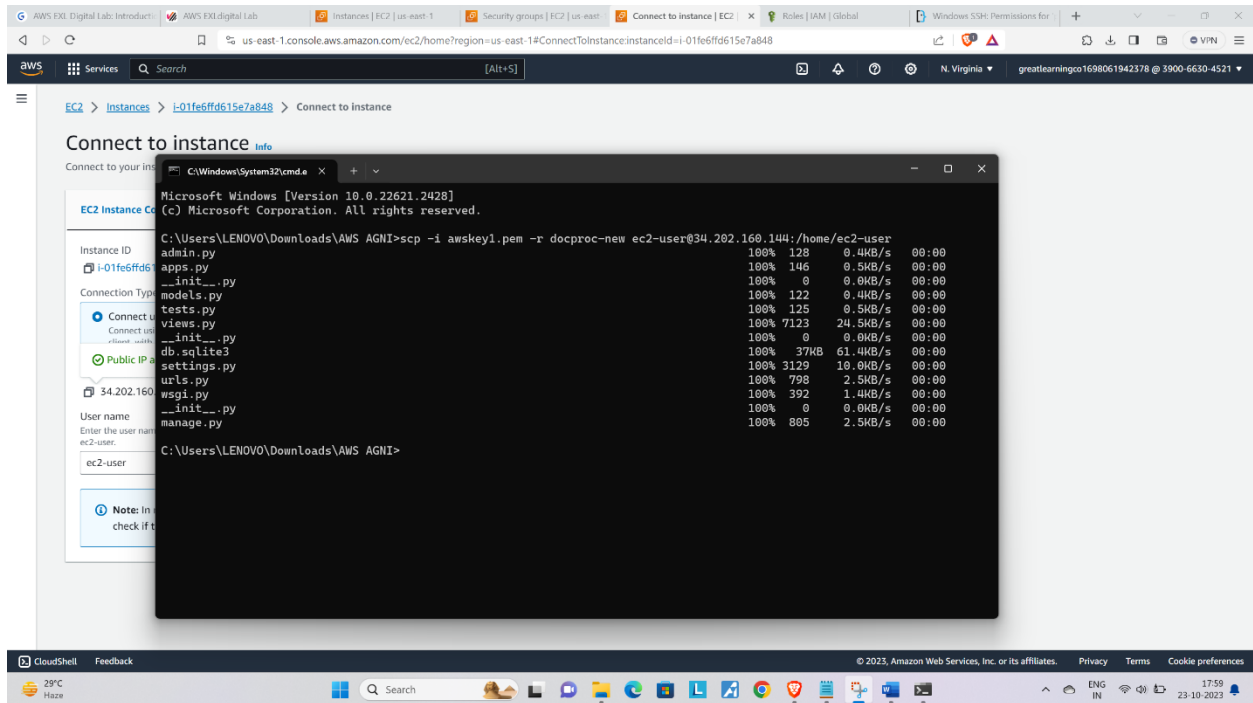
Expected screens
hots

- 1) Modifying of the [views.py](#) file to point to the target bucket
- 2) Copying the folder to the EC2 instance

<Insert screenshot for c(1) here>

<Insert screenshot for c(2) here>





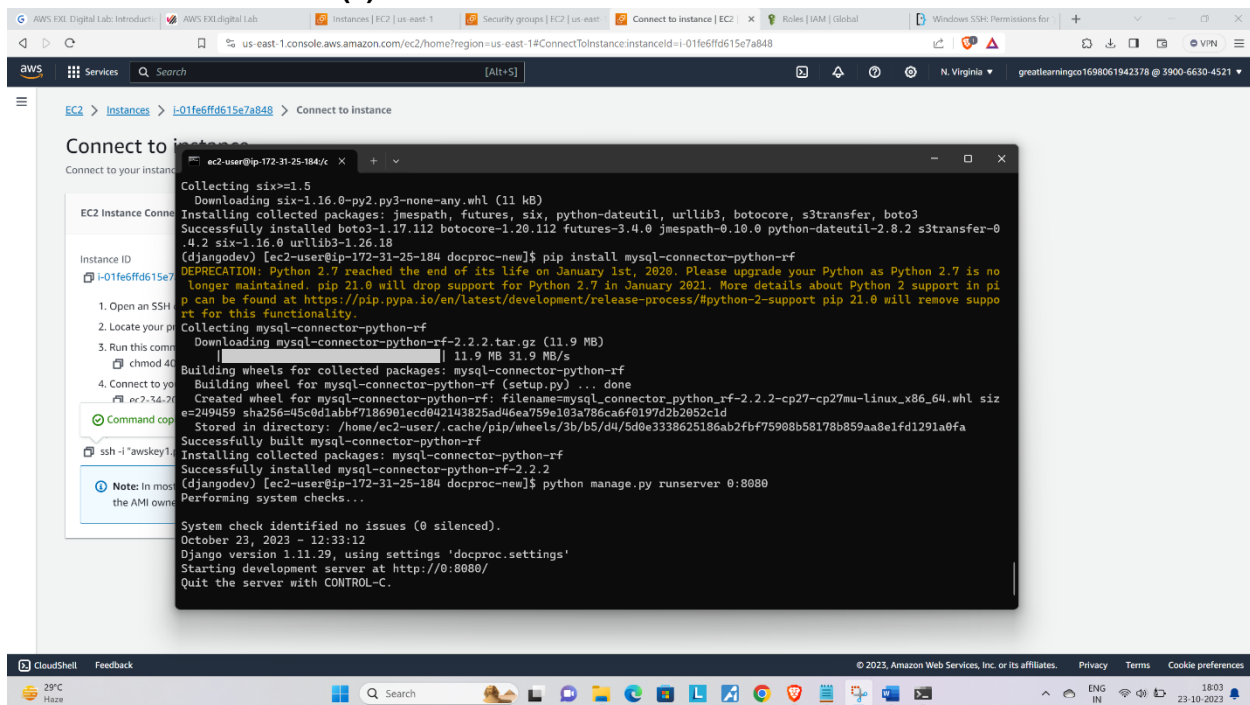
Step 3: Creation and Verification of SNS subscription and Generation of CSV file

Step number a

Step name Starting the
EC2 custom
program

Instructions	1) Log into the EC2 instance using SSH 2) Run the following commands after successful SSH to start the server <pre> sudo cp -r docproc-new /opt sudo chown ec2-user:ec2-user -R /opt cd /opt/docproc-new sudo yum update sudo yum install python-pip -y python -m pip install --upgrade pip setuptools sudo pip install virtualenv virtualenv ~/.virtualenvs/djangodev source ~/.virtualenvs/djangodev/bin/activate pip install django pip install boto3 pip install mysql-connector-python-rf python manage.py runserver 0:8080 </pre> <p>Keep this terminal window open throughout the rest of the exercise</p>
Expected screenshots	1) Server in waiting state

<Insert screenshot for a(1) here>

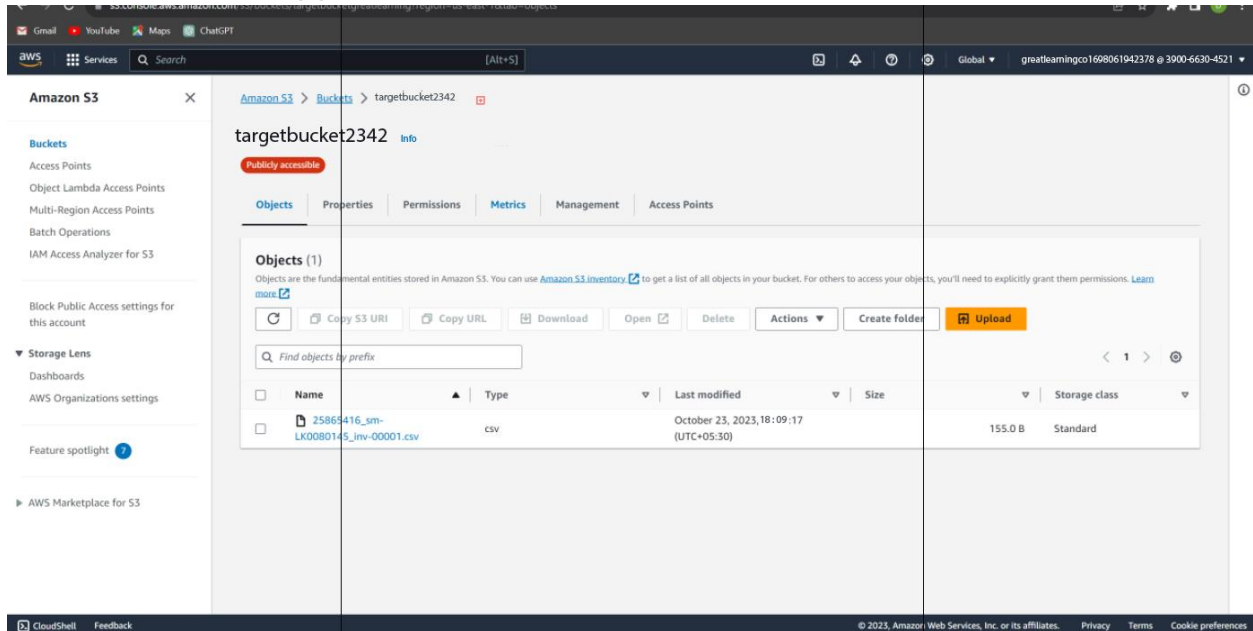


Step number	b
Step name	Creation of SNS subscription
Instructions	<p>1) Navigate to SNS in the AWS Console and select the topic S3ToEC2Topic</p> <p>2) Click on Create Subscription</p> <p>3) Enter the following details</p> <p>Protocol : HTTP</p> <p>Endpoint : http://<host>:8080/sns where <host> in the public IP of the EC2 instance</p> <p>Click on Create Subscription</p> <p>4) In the EC2 terminal window, look for the field "SubscribeURL" and copy the entire link given</p> <p>Note: If a message is seen "ValueError: No JSON object could be decoded", it can be safely ignored</p> <p>5) Paste that link into a browser window to verify the SNS subscription (Ignore any messages received in the web browser)</p>
Expected screenshots	<p>1)</p> <p>Subscription URL in EC2 terminal Window</p>

<Insert screenshot for b(1) here>

Expected screenshots 1) Generated CSV file in the target S3 bucket

<Insert screenshot c(1) here>



Answer the following questions

Q1 Which of the following properties of an AWS resource is sufficient and necessary to uniquely identify it across all of AWS?

- a) ARN
- b) Region and ARN
- c) ARN and Account number
- d) Depends on the resource used

Enter your answer here

a

Q2 Which of the following step numbers in Step 1 allowed S3 to publish to the SNS topic created?

- a) 1(a)
- b) 1(c)
- c) 1(d)
- d) 1(b)

Enter your answer here

b

Q3 Which port is being used by SNS to send the notification to the custom program?

a) 8081

b) 80

c) 8080

d) 8065

Enter your answer here

c

Q4 How many IAM roles can be attached to an EC2 instance at a time?

a) 2

b) 3

c) 1

d) Depends on the policies required

Enter your answer here

c

Q5 As a product manager, how would you describe the benefits of this architecture to an client, as compared to an equivalent on-premises architecture?

In a cloud environment, a third-party service provider hosts all of these resources for us. It's a really simple and user-friendly setup. This architecture's main qualities are its simple design and general accessibility. Instant provisioning eliminates the need for installation and configuration, allowing users to access the program immediately.

Grades distribution	
MCQs	10 (2.5 mark each)
Subjective questions	6 marks
Implementation screenshots	24 marks (2 marks each)
Total	40 marks